

Spam_email_detect

October 11, 2023

```
[48]: import numpy as np
import pandas as pd
import string
import matplotlib.pyplot as plt
import seaborn as sns
import nltk
from nltk.corpus import stopwords
from collections import Counter
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, log_loss, \
    precision_score, classification_report, f1_score
```

```
[49]: # Read the dataset full of spams and non-spams.
dataframe = pd.read_csv('spam.csv', encoding = 'latin1')
print(dataframe)
```

	v1	v2	Unnamed: 2	\
0	ham	Go until jurong point, crazy.. Available only ...	NaN	
1	ham	Ok lar... Joking wif u oni...	NaN	
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	
3	ham	U dun say so early hor... U c already then say...	NaN	
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN	
...
5567	spam	This is the 2nd time we have tried 2 contact u...	NaN	
5568	ham	Will I_ b going to esplanade fr home?	NaN	
5569	ham	Pity, * was in mood for that. So...any other s...	NaN	
5570	ham	The guy did some bitching but I acted like i'd...	NaN	
5571	ham	Rofl. Its true to its name	NaN	
	Unnamed: 3	Unnamed: 4		
0	NaN	NaN		
1	NaN	NaN		
2	NaN	NaN		
3	NaN	NaN		
4	NaN	NaN		

```
...      ...      ...
5567      NaN      NaN
5568      NaN      NaN
5569      NaN      NaN
5570      NaN      NaN
5571      NaN      NaN
```

[5572 rows x 5 columns]

```
[50]: # Let's drop the unnecessary columns and extract the data we want, which is
      ↪ emails with some kind of content i.e not empty
      dataframe = dataframe.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'],axis=1)
      dataframe.head()
```

```
[50]:      v1      v2
0   ham  Go until jurong point, crazy.. Available only ...
1   ham                Ok lar... Joking wif u oni...
2  spam  Free entry in 2 a wkly comp to win FA Cup fina...
3   ham  U dun say so early hor... U c already then say...
4   ham  Nah I don't think he goes to usf, he lives aro...
```

```
[51]: # Let's also change 'v1' and 'v2' names for clearer set
      dataframe.columns = ['Label', 'E-mail']
```

```
[52]: dataframe.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Label   5572 non-null       object
1   E-mail  5572 non-null       object
dtypes: object(2)
memory usage: 87.2+ KB
```

```
[53]: # Let's check and get rid of duplicates if there are any
      dataframe.duplicated().sum()
```

```
[53]: 403
```

```
[54]: dataframe = dataframe.drop_duplicates()
      dataframe.shape
```

```
[54]: (5169, 2)
```

```
[55]: # Let's see how many spam- and non-spam messages we have (datapoints)
count = dataframe['Label'].value_counts()
```

```
[56]: count_of_non_spams = count['ham']
count_of_spams = count['spam']

print("Count of non-spams:", count_of_non_spams)
print("Count of spams:", count_of_spams)
```

Count of non-spams: 4516

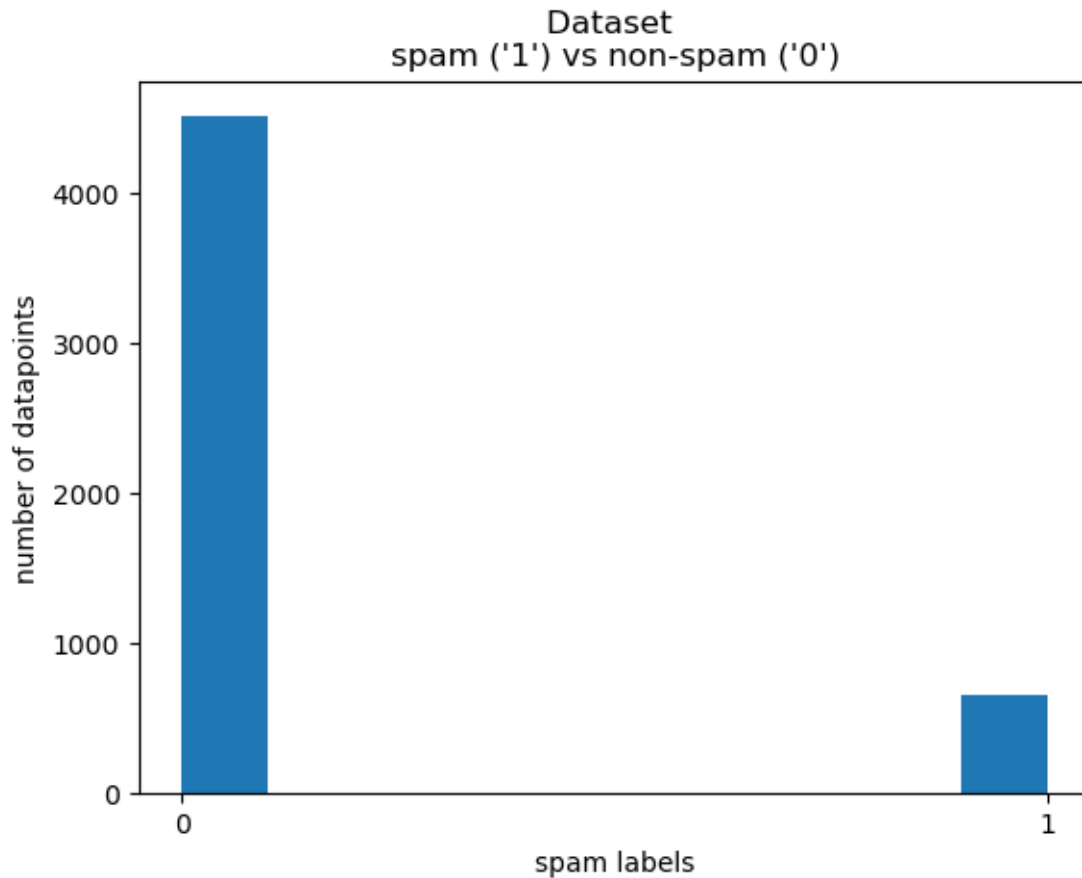
Count of spams: 653

```
[57]: # Let's change 'ham' and 'spam' into binary values, where 1 = spam and 0 =
      ↪ non-spam
dataframe.loc[dataframe['Label'] == 'spam', 'Label',] = 1
dataframe.loc[dataframe['Label'] == 'ham', 'Label',] = 0
```

```
[58]: dataframe.head()
```

```
[58]:   Label                                     E-mail
0      0  Go until jurong point, crazy.. Available only ...
1      0                                     Ok lar... Joking wif u oni...
2      1  Free entry in 2 a wkly comp to win FA Cup fina...
3      0  U dun say so early hor... U c already then say...
4      0  Nah I don't think he goes to usf, he lives aro...
```

```
[59]: fig, ax = plt.subplots()
ax.hist(dataframe['Label'])
ax.set_title('Dataset \nspam (\n'1\n') vs non-spam (\n'0\n')')
ax.set_xlabel("spam labels")
ax.set_ylabel('number of datapoints')
ax.set_xticks([0,1])
plt.show()
```



```
[60]: def remove_special_characters(text):
      # Remove punctuation using string.punctuation and str.translate
      translator = str.maketrans('', '', string.punctuation)
      text = text.translate(translator)
      return text.lower()
```

```
[61]: dataframe['E-mail'] = dataframe['E-mail'].apply(remove_special_characters)
```

```
[62]: dataframe.head()
```

```
[62]:  Label                                E-mail
0      0  go until jurong point crazy available only in ...
1      0                                ok lar joking wif u oni
2      1  free entry in 2 a wkly comp to win fa cup fina...
3      0          u dun say so early hor u c already then say
4      0  nah i dont think he goes to usf he lives aroun...
```

```
[63]: # Let's separate the columns of the data into their own parameters:
      X = dataframe['E-mail']
```

```
y = dataframe['Label']
```

```
[64]: print(X, y)
```

```
0      go until jurong point crazy available only in ...
1                      ok lar joking wif u oni
2      free entry in 2 a wkly comp to win fa cup fina...
3          u dun say so early hor u c already then say
4      nah i dont think he goes to usf he lives aroun...

...
5567    this is the 2nd time we have tried 2 contact u...
5568                      will i b going to esplanade fr home
5569    pity was in mood for that soany other suggest...
5570    the guy did some bitching but i acted like id ...
5571                      rofl its true to its name
Name: E-mail, Length: 5169, dtype: object 0      0
1      0
2      1
3      0
4      0
..
5567    1
5568    0
5569    0
5570    0
5571    0
Name: Label, Length: 5169, dtype: object
```

```
[65]: # Let's make our sets. Training, valid and test with 7:1,5:1,5
# Split the data into "train" (70%) and "temporary" (30%)
X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.3,
↳random_state=10)

# Split the "temp" set into "validation" (50%) and "test" (50%)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5,
↳random_state=10)
```

```
[66]: # Let's check the shapes of the original-, test-, val- and train data's X and y
↳to see that we have succeeded in the grouping
print(X.shape)
print(X_train.shape)
print(X_val.shape)
print(X_test.shape)

(5169,)
(3618,)
(775,)
```

(776,)

```
[67]: print(y.shape)
      print(y_train.shape)
      print(y_val.shape)
      print(y_test.shape)
```

(5169,)

(3618,)

(775,)

(776,)

```
[68]: # Since the lengths match, the grouping and distribution has succeeded
```

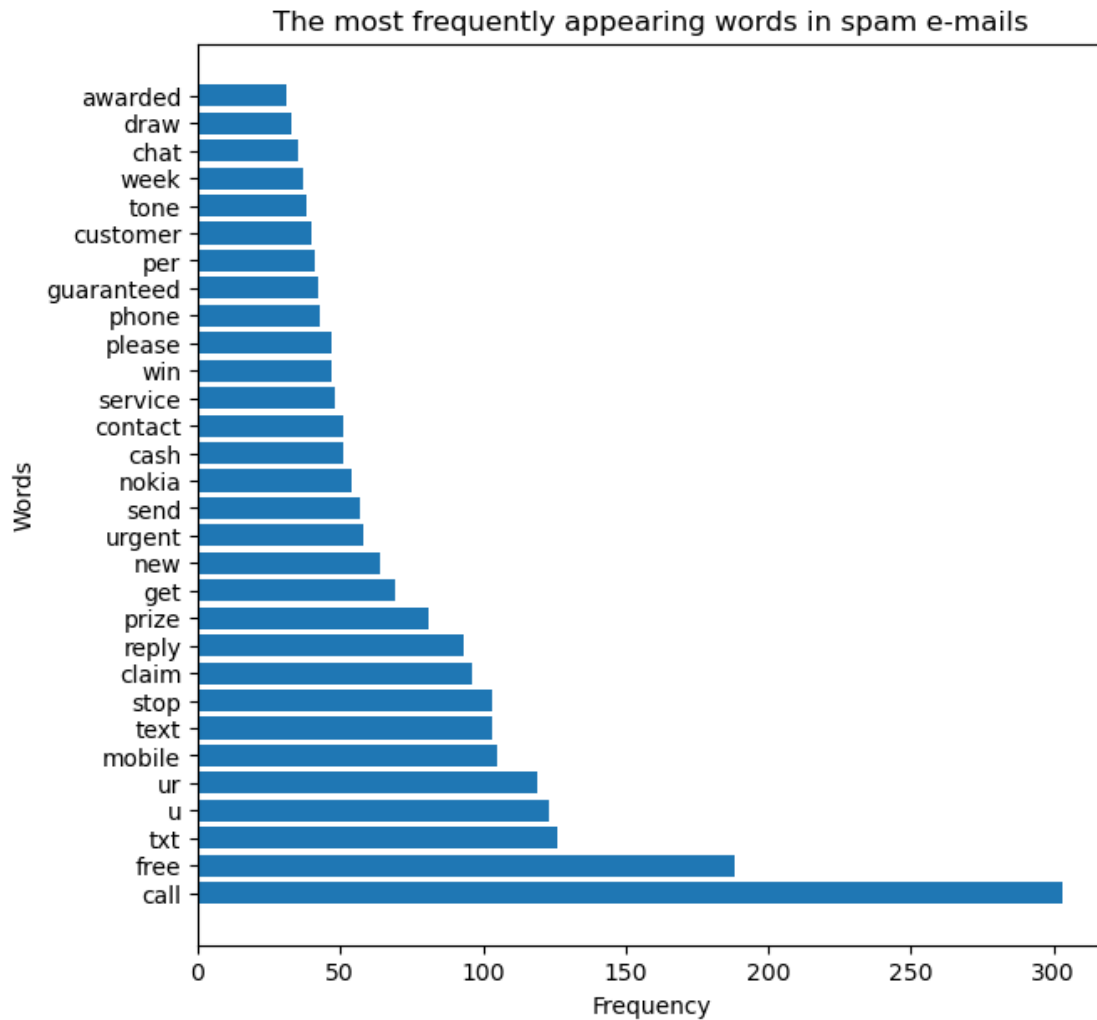
```
[69]: # Let's now extract the features from data's. This is where TfidfVectorizer
      ↪ (Term Frequency - Inverse Document Frequency) comes in.
      # This allows us to assess the quantity and quality of words in emails and help
      ↪ us determine the spams.
      # min_df ensures, that extremely infrequent words are cut-off of our equation,
      ↪ which makes these words more unique, making them less likely a spam
      # Frequently appearing words can be assumed to be part of a spam message
      features = TfidfVectorizer(min_df = 1, stop_words = 'english', lowercase = True)
      X_train = features.fit_transform(X_train)
      X_test = features.transform(X_test)
      X_val = features.transform(X_val)

      y_val = y_val.astype(int)
      y_train = y_train.astype(int)
      y_test = y_test.astype(int)
```

```
[70]: # Let's see what kind of words appear most frequently in spams
      spam_most_used_words = []
      stop_words = set(stopwords.words('english'))
      for sentence in dataframe[dataframe['Label'] == 1]['E-mail'].tolist():
          words = sentence.split()
          filtered_words = [word for word in words if word.lower() not in stop_words
          ↪ and word.isalpha()]
          spam_most_used_words.extend(filtered_words)

      filter_df = pd.DataFrame(Counter(spam_most_used_words).most_common(30))
```

```
[71]: plt.figure(figsize=(7, 7))
      plt.barh(filter_df[0], filter_df[1])
      plt.xlabel('Frequency')
      plt.ylabel('Words')
      plt.title('The most frequently appearing words in spam e-mails')
      plt.show()
```



```
[72]: print(X_train)
```

```
(0, 5190)    0.4521241713824104
(0, 5793)    0.3721282557866962
(0, 3353)    0.1896790597744401
(0, 1766)    0.4152584200119519
(0, 1545)    0.4152584200119519
(0, 4899)    0.32458962544662484
(0, 3836)    0.22214069331416178
(0, 5930)    0.34862892281720176
(1, 6762)    0.2927895555121226
(1, 2141)    0.2927895555121226
(1, 3704)    0.2927895555121226
(1, 7021)    0.2549505306455541
(1, 5676)    0.2927895555121226
(1, 5574)    0.2927895555121226
```

(1, 6834)	0.2927895555121226
(1, 1213)	0.27882429728000246
(1, 2895)	0.2927895555121226
(1, 5482)	0.2927895555121226
(1, 1384)	0.2927895555121226
(1, 7037)	0.2927895555121226
(2, 1470)	0.4966035959316888
(2, 3151)	0.4652464853142722
(2, 3024)	0.44545358337764496
(2, 7095)	0.4874947093614859
(2, 2248)	0.3175698187973922
:	:
(3614, 3237)	0.47750925337263866
(3614, 3737)	0.5333301936875326
(3614, 6544)	0.4117250459123929
(3614, 1104)	0.43184393400337695
(3614, 6514)	0.36268046635799
(3615, 6512)	0.3956700211092583
(3615, 4978)	0.3767976470253049
(3615, 6967)	0.30188635690203547
(3615, 1790)	0.31632251851959775
(3615, 7230)	0.33736016680322134
(3615, 1355)	0.32075873098598895
(3615, 5510)	0.24584744086271954
(3615, 2896)	0.29340019975909737
(3615, 1036)	0.23814357106780987
(3615, 6372)	0.23509034545306035
(3615, 4609)	0.18727237215508966
(3616, 1999)	0.664184765345535
(3616, 5468)	0.578348015438433
(3616, 6893)	0.3392430085479584
(3616, 3973)	0.3305848630422339
(3617, 2593)	0.5759501518340481
(3617, 2487)	0.5013270511240697
(3617, 7032)	0.4202853532449395
(3617, 2992)	0.3669515613171994
(3617, 2039)	0.3250528940160777

```
[73]: def generate_confusion_matrix(y_true, y_pred, title_main):
    # visualize the confusion matrix
    ax = plt.subplot()
    c_mat = confusion_matrix(y_true, y_pred)
    sns.heatmap(c_mat, annot=True, fmt='g', ax=ax)

    ax.set_xlabel('Predicted labels', fontsize=15)
    ax.set_ylabel('True labels', fontsize=15)
    ax.set_title(title_main, fontsize=15)
```



```
[74]: # Let's train our model and then do validation
      clf_1 = LogisticRegression(solver = 'liblinear', class_weight = 'balanced').
      ↪fit(X_train, y_train)
```

```
[75]: y_pred_train_lr = clf_1.predict(X_train)
      tr_error = log_loss(y_train,y_pred_train_lr)
      accuracy = accuracy_score(y_train, y_pred_train_lr)
      precision = precision_score(y_train, y_pred_train_lr)
```


```
[76]: print("Error:",tr_error)
      print("Accuracy:",accuracy)
      print("Precision:",precision)
```

Error: 0.2590201736089129

Accuracy: 0.9928137092316197

Precision: 0.9491525423728814

```
[77]: generate_confusion_matrix(y_train, y_pred_train_lr, 'Logistic Regression_
      ↪\nConfusion matrix (training)' )
```




```
[78]: # Let's use our validation set
y_pred_val_lr = clf_1.predict(X_val)
val_error = log_loss(y_val ,y_pred_val_lr)
accuracy = accuracy_score(y_val, y_pred_val_lr)
precision = precision_score(y_val, y_pred_val_lr)
```

```
[79]: print("Error:",val_error)
print("Accuracy:",accuracy)
print("Precision:",precision)
```

Error: 1.3022223159939101
Accuracy: 0.9638709677419355
Precision: 0.9080459770114943

```
[80]: generate_confusion_matrix(y_val, y_pred_val_lr, 'Logistic Regression_1')
↪\nConfusion matrix (validation)'
```

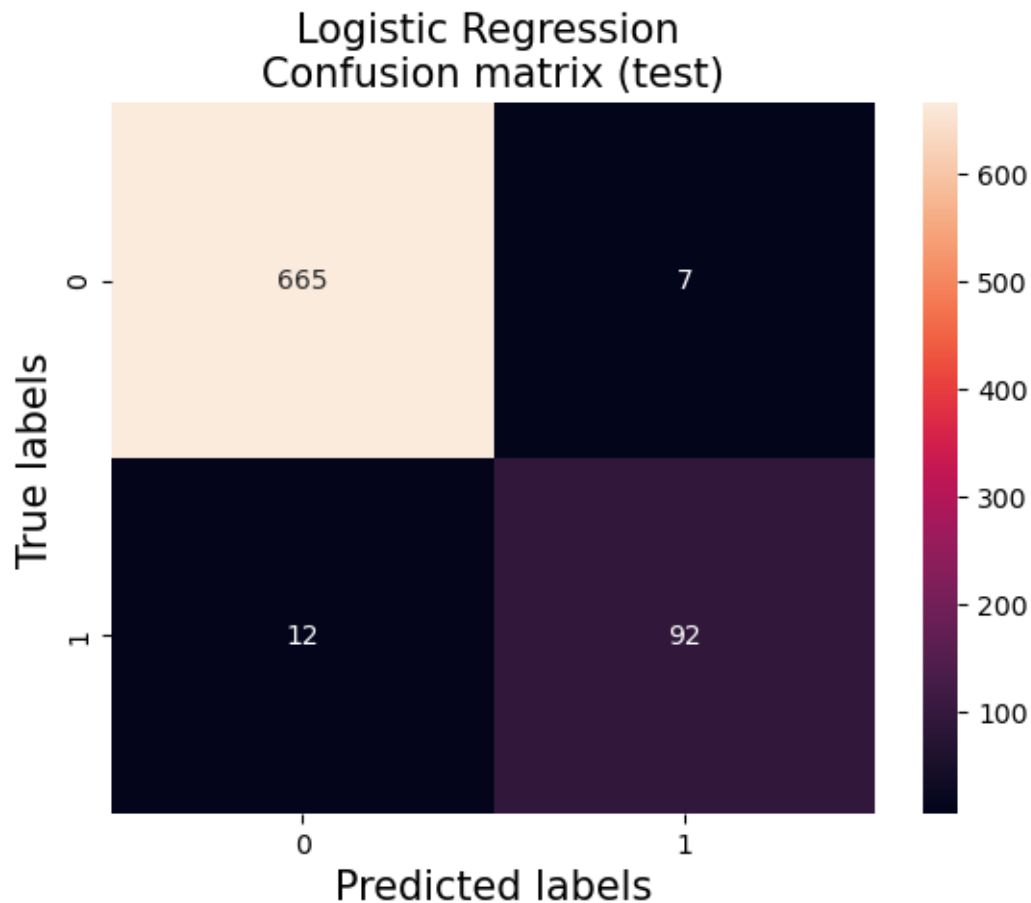


```
[81]: # Let's also use our test set
y_pred_tst_lr = clf_1.predict(X_test)
tst_error = log_loss(y_test,y_pred_tst_lr)
accuracy = accuracy_score(y_test, y_pred_tst_lr)
precision = precision_score(y_test, y_pred_tst_lr)
```

```
[82]: print("Error:",tst_error)
print("Accuracy:",accuracy)
print("Precision:",precision)
```

Error: 0.8825121319500335
Accuracy: 0.9755154639175257
Precision: 0.9292929292929293

```
[83]: generate_confusion_matrix(y_test, y_pred_tst_lr, 'Logistic Regression_
↳\nConfusion matrix (test)')
```



```
[84]: # Let's see how well can we do with another ML-method, such as random forest
      ↪ classifier
      # Let's find the optimal parameters to best estimation
      n_estimators = [10, 50, 100, 200]
      max_depth = [None, 2, 3, 5]

      # Create a parameter grid: map the parameter names to the values that should be
      ↪ searched
      param_grid = {'n_estimators': n_estimators, 'max_depth': max_depth}

      # Instantiate the grid
      grid = GridSearchCV(RandomForestClassifier(), param_grid, cv=5)

      # Fit the grid with data
      grid.fit(X_train, y_train)

      # Examine the best model
      print(grid.best_score_)
      print(grid.best_params_)
```

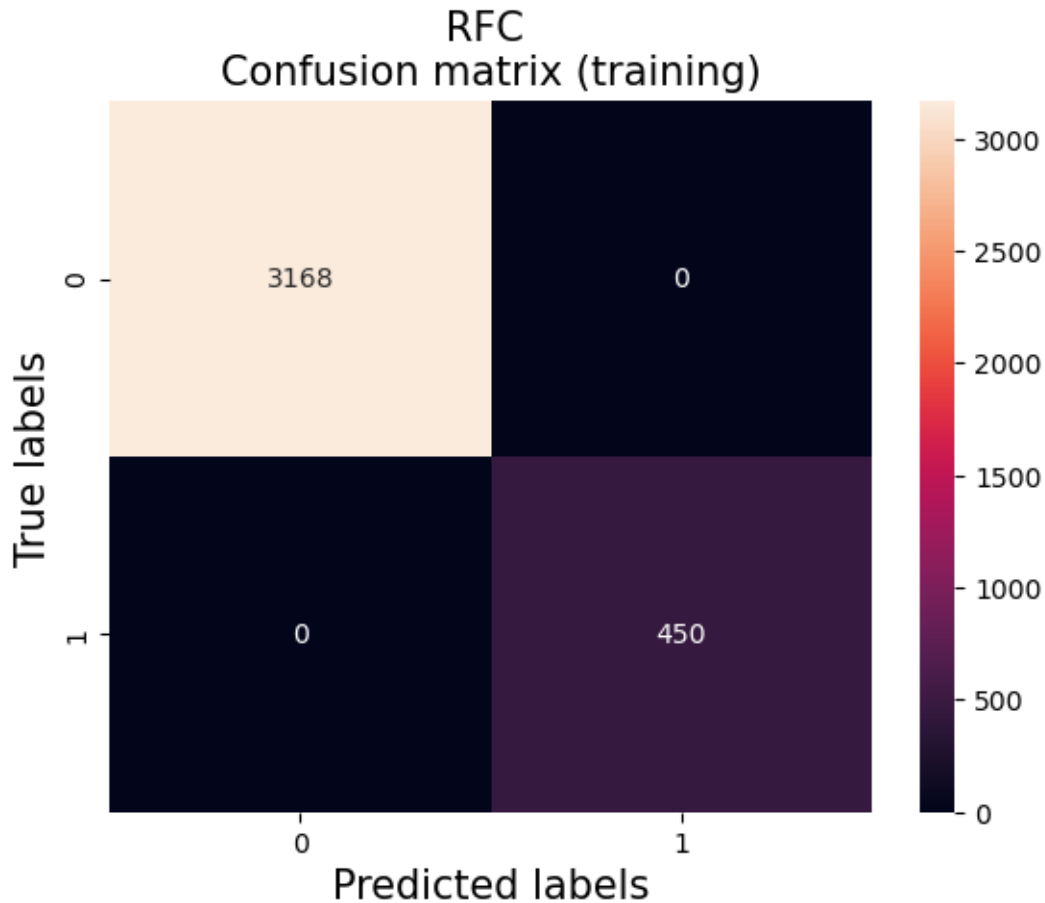
```
0.9690443440850356
{'max_depth': None, 'n_estimators': 100}
100
```

```
[85]: # Now let's build our model
      rfc = RandomForestClassifier(n_estimators = grid.best_params_['n_estimators'],
      ↪ random_state = 2, criterion = 'gini', max_depth = grid.
      ↪ best_params_['max_depth'], class_weight = 'balanced')
      # Train our model and see the performance
      rfc.fit(X_train, y_train)
      y_pred_train_rfc = rfc.predict(X_train)
      y_accuracy_train = accuracy_score(y_train, y_pred_train_rfc)
      precision_train = precision_score(y_train, y_pred_train_rfc)
```

```
[86]: print("Accuracy: ", y_accuracy_train)
      print("Precision: ", precision_train)
```

```
Accuracy:  1.0
Precision:  1.0
```

```
[87]: generate_confusion_matrix(y_train, y_pred_train_rfc, 'RFC \nConfusion matrix
      ↪ (training)')
```

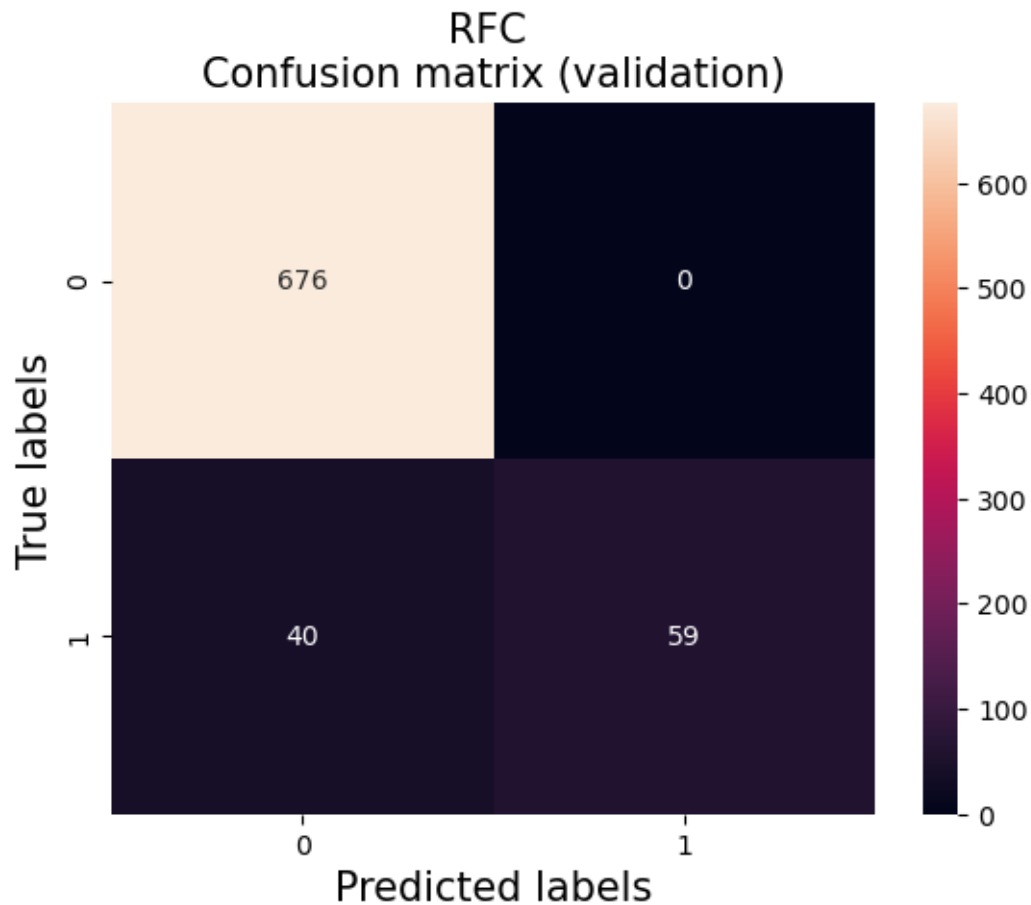


```
[88]: # How well does the validation set perform
y_pred_val_rfc = rfc.predict(X_val)
y_accuracy_val = accuracy_score(y_val, y_pred_val_rfc)
precision_val = precision_score(y_val, y_pred_val_rfc)
```

```
[89]: print("Accuracy: ", y_accuracy_val)
print("Precision: ", precision_val)
```

```
Accuracy:  0.9483870967741935
Precision:  1.0
```

```
[90]: generate_confusion_matrix(y_val, y_pred_val_rfc, 'RFC \nConfusion matrix_\n↪(validation)')
```




```
[91]: y_pred_tst_rfc = rfc.predict(X_test)
      y_accuracy_tst = accuracy_score(y_test, y_pred_tst_rfc)
      precision_tst = precision_score(y_test, y_pred_tst_rfc)
```

```
[92]: print("Accuracy: ", y_accuracy_tst)
      print("Precision: ", precision_tst)
```

```
Accuracy:  0.9548969072164949
Precision:  1.0
```

```
[93]: generate_confusion_matrix(y_test, y_pred_tst_rfc, 'RFC \nConfusion matrix_\n↳(test)')
```



```
[94]: print("\033[1m" + "Logistic Regression classification report of training\ndata\n", "\033[0m" + classification_report(y_train, y_pred_train_lr))
print("\033[1m" + "Logistic Regression classification report of validation\ndata\n", "\033[0m" + classification_report(y_val, y_pred_val_lr))
print("\033[1m" + "Logistic Regression classification report of test data\n", "\033[0m" + classification_report(y_test, y_pred_tst_lr))
print("-----")
print("\033[1m" + "RFC classification report for training data\n", "\033[0m" + classification_report(y_train, y_pred_train_rfc))
print("\033[1m" + "RFC classification report of validation data\n", "\033[0m" + classification_report(y_val, y_pred_val_rfc))
print("\033[1m" + "RFC classification report of test data\n", "\033[0m" + classification_report(y_test, y_pred_tst_rfc))
```

Logistic Regression classification report of training data

	precision	recall	f1-score	support
0	1.00	0.99	1.00	3168

1	0.95	1.00	0.97	450
accuracy			0.99	3618
macro avg	0.97	0.99	0.98	3618
weighted avg	0.99	0.99	0.99	3618

Logistic Regression classification report of validation data

	precision	recall	f1-score	support
0	0.97	0.99	0.98	676
1	0.91	0.80	0.85	99
accuracy			0.96	775
macro avg	0.94	0.89	0.91	775
weighted avg	0.96	0.96	0.96	775

Logistic Regression classification report of test data

	precision	recall	f1-score	support
0	0.98	0.99	0.99	672
1	0.93	0.88	0.91	104
accuracy			0.98	776
macro avg	0.96	0.94	0.95	776
weighted avg	0.98	0.98	0.98	776

RFC classification report for training data

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3168
1	1.00	1.00	1.00	450
accuracy			1.00	3618
macro avg	1.00	1.00	1.00	3618
weighted avg	1.00	1.00	1.00	3618

RCF classification report of validation data

	precision	recall	f1-score	support
0	0.94	1.00	0.97	676
1	1.00	0.60	0.75	99
accuracy			0.95	775
macro avg	0.97	0.80	0.86	775

weighted avg	0.95	0.95	0.94	775
--------------	------	------	------	-----

RCF classification report of test data

	precision	recall	f1-score	support
0	0.95	1.00	0.97	672
1	1.00	0.66	0.80	104
accuracy			0.95	776
macro avg	0.98	0.83	0.89	776
weighted avg	0.96	0.95	0.95	776

[]: